

Introduction to CouchDB

Sophia Sandhu



CouchDB
relax

<http://couchdb.apache.org>



Agenda

- CouchDB History
- CouchDB Characteristics
- Who is using Couch DB
- CouchDB Examples
- CouchDB
 - Documents
 - Storage
 - Views
 - Compaction
 - Replication
- CouchDB Lab





CouchDB's History

- Development started around 2005 by Damien Katz (Lotus Notes, IBM, MySQL)
- Initially written in C++, used XML and a query language similar to Formula (Lotus Notes)
- 2007 - adopted Erlang, JSON and MapReduce with JavaScript, dropping C++, XML and the custom query language
- Drew interest from IBM and got sponsored by IBM
- 2008 - became an Incubator project of the Apache Software Foundation
- 2009 - became a top level Apache project (alongside with httpd, Tomcat, etc)
- 2010 - Proposal to the HTML5 Indexed DB API to make it possible for someone to implement a CouchDB API on top of it (<http://www.w3.org/TR/IndexedDB>)



CouchDB's Characteristics

- It's a document-oriented DBMS (i.e. Not a relational DBMS)
- Data is modeled as “documents”, resembling real world documents and promoting self-contained data
- A document is a JSON structure containing any kind and number of fields
- Arbitrary binary data (video, audio, images, etc) is supported as document attachments
- Communication with the outside world done exclusively through an HTTP RESTful API
- Has incremental peer to peer replication -very simple and it uses exclusively the public HTTP RESTful API (no hidden/special APIs or protocols)



CouchDB's Characteristics

- Uses MapReduce for computing views
- ACID properties guaranteed at a single document level
- Doesn't support transactions for updating/adding/removing multiple documents at once -a choice by design
- Instead of locks, it uses Multi Version Concurrency Control (MVCC), i.e. revision numbers, to manage concurrent requests
- Main goals: availability and partition tolerance



CouchDB' vs RDBMS

- RDBMSs do not scale so gracefully due to their focus on distributed consistency:
- Before writing to a node, synchronization protocols must run between all the nodes and then some distributed commit protocol takes action.
- Client requests are blocked until all nodes reach an agreement and data is committed to all nodes .
- More nodes => more communication overhead => more time for reaching agreements

Who is using Couch DB?



- - Canonical Ltd, Ubuntu One service
- BBC UK
- Mozilla is using it for Raindrop.
- Facebook
- Assay Depot, a big marketplace in USA for pharmaceutical research services.

HTTP RESTful API examples



```
$ curl -X PUT http://localhost:5984/albums
{"ok":true}
```

```
$ curl -X PUT http://localhost:5984/albums/album1 -d @-
{ "artist": "Megadeth",
  "title": "Endgame",
  "year": 2009 }
{"ok":true,"id":"album1","rev":"1-35cc06dfb32b482b2e9031a4e3ebaef0"}
```

```
$ curl -X GET http://localhost:5984/albums/album1
{
  "_id":"album1", "_rev":"1-35cc06dfb32b482b2e9031a4e3ebaef0", "artist":"Megadeth",
  "title":"Endgame",
  "year":2009
}
```


HTTP RESTful API examples



To update a document, we need to supply its latest revision, otherwise we get an HTTP response with error code 409

```
$ curl -X PUT http://localhost:5984/albums/album1 -d @-{"artist": "Megadeth",  
"title": "Endgame",  
"year": 2010  
}  
{"error": "conflict", "reason": "Document update conflict."}
```

An attribute named “_rev” specifies the target revision:

```
$ curl -X PUT http://localhost:5984/albums/album1 -d @-{"_rev": "1-35cc06dfb32b482b2e9031a4e3ebaef0",  
"artist": "Megadeth",  
"title": "Endgame",  
"year": 2010  
}
```

HTTP RESTful API examples



```
$ curl -X DELETE http://localhost:5984/albums/album1?rev=\2-a568258ebba8a79c2ebe9fb5c3700ac2
```

```
{"ok":true,"id":"album1","rev":"3-ea701815806524a3bdf7c81e35276801"}
```

```
$ curl -X GET http://localhost:5984/albums/album1  
{"error":"not_found","reason":"deleted"}
```

```
$ curl -X GET http://localhost:5984/albums/album1?rev=\3ea701815806524a3bdf7c81e35276801  
{"_id":"album1","_rev":"3-ea701815806524a3bdf7c81e35276801","_deleted":true}
```

```
$ curl -X POST http://localhost:5984/albums/_purge \-d  
'{ "album1": ["3-ea701815806524a3bdf7c81e35276801"] } '  
{"purge_seq":1,"purged":{"album1":["3-ea701815806524a3bdf7c81e35276801"]}}
```

Document Attachments



- Special attribute “_attachments” in a document with attachments:

```
$ curl -X GET http://localhost:5984/albums/album1
```

```
{  
  "_id": "album1",  
  "_rev": "2-2e0e230f5a3c44aa1dac72e5506ae213",  
  "artist": "Megadeth",  
  "title": "Endgame",  
  "year": 2010,  
  "_attachments":  
    { "cover.png":  
      { "content_type": "image/png",  
        "length": 2157,  
        "stub": true } } }
```

CouchDB Storage



- Each DB corresponds to a single file
- All data is written in file append mode – existing data is never ever overwritten => no need for data integrity check and recovery on server startup
- Uses B-Trees – a data structure for efficient operations on very large data sets stored in disk. For each tree operation (search, insert, delete) only 2 nodes at most need to be in RAM memory for each step
- Each DB file contains 2 B-Trees:
 - by_doc_id_btree – keys are document IDs, values contain pointers (offsets) to where the document is located within the DB file and other metadata
 - by_db_seq_btree – keys are DB sequence numbers, values are information about what happened (document added or updated) for a particular DB sequence number

CouchDB Views



- To some extent, they have the same end result as SQL in RDBMSs.
- Views can be defined using a Map and (optionally) a Reduce function. JavaScript is the preferred language, but there is 3rd party support for others (except for Erlang which is natively supported as well)
- Views are mapped to individual files as well => they can be stored on different disks for higher IO performance.
- Like a DB file, a view file contains a B-Tree (a single one)
- View B-Tree keys are the ones emitted by a Map function.
- Values emitted by a Map function are stored in the leaf nodes of the B-Tree.
- A non-leaf node of the B-Tree stores the result of a Reduce operation over the values of its children => Reduce functions are used as well for doing a “Re-Reduce”

CouchDB Views



- When a document is updated/added/deleted, Map and Reduce computations are minimized:
 - The Map value for the affected document is re-calculated
 - The Reduce value of the parent node is re-computed
 - Finally re-computation of Reduce values takes place all the way up the B-Tree(Re-Reduce operations)
- Typically the B-Tree has from 4 to 6 levels for a DB with millions of documents.
- Updating the view after a few documents were updated and/or added is cheap.
- Map and Reduce functions are executed by the View Server (it uses Mozilla's SpiderMonkey JavaScript engine)

Couch View Examples



Example Dataset

```
{
  "_id": "album1",
  "artist": "Megadeth",
  "title": "Endgame",
  "year": 2010
}

{
  "_id": "album2",
  "artist": "Slayer",
  "title": "World Painted Blood",
  "year": 2009
}

{
  "_id": "album3",
  "artist": "Arcturus",
  "title": "Sideshow Symphonies",
  "year": 2005
},
```

```
{
  "_id": "album4",
  "artist": "Pantera",
  "title": "Reinventing the Steel",
  "year": 2009
}

{
  "_id": "album5",
  "artist": "Slayer",
  "title": "South of Heaven",
  "year": 2009
}
```

Couch View Examples



- View definitions (Map and Reduce functions) are stored in “design” documents – their ID is prefixed with “_design/” and have a special meaning in CouchDB

```
{
  "_id": "_design/foobar",
  "language": "javascript",
  "views": {
    "by_year": {
      "map": "function(doc) {
                if (doc.year) {
                  emit(doc.year, 1);
                }
              }",
      "reduce": "function(keys, values, rereduce) {
                  return sum(values);
                }"
    }
  }
}
```


Couch View Examples



Getting the Reduce result for all emitted values, ignoring key values:

```
$ curl http://localhost:5984/albums/_design/foobar/_view/by_year
{
  "update_seq": 6,
  "rows": [
    {"key": null, "value": 5}
  ]
}
```

Getting the Reduce for all emitted values but for each distinct key:

```
$ curl http://localhost:5984/albums/_design/foobar/_view/by_year?group=true
{
  "update_seq": 6,
  "rows": [
    {"key": 2005, "value": 1},
    {"key": 2009, "value": 3},
    {"key": 2010, "value": 1}
  ]
}
```

Compaction



- DB files (as well as View files) are written in append mode only, not overriding any previously written data
- DB files and View files will grow indefinitely!!! Unless no more documents are added or updated...
- A DB or View compaction operation can be triggered.
- A compaction operation does the following
 - Creates a new file
 - Starts traversing the DB or View B-Tree and lookups the most recent data pointed by each node (documents or map or reduce values)
 - Writes that most recent data to the new file
 - When the original DB or View file is not being accessed to serve a request, it deletes it and renames the compacted file to the original DB/View file name.

Replication



- Peer to peer replication, no concept of master or slave roles.
- A replication specifies a source and a target DB (unidirectional)
- Any of the DBs, or both, can be local or remote
- Requesting a replication is very simple

```
$ curl -X POST http://localhost:5984/_replicate  
-d @-
```

```
{ "source": "http://example.com/somedb",  
  "target": "somedb_copy" }
```

Replication



- It creates a checkpoint history document (a document with the prefix “_local/” as part of its ID) in the source and target DBs
- It uses the `_changes` API from the source DB
 - For each received line from `_changes`, which mentions a source DB sequence number `SEQ_NUM`, retrieves the corresponding Document from the source DB2
 - Inserts the document into the target DB (if it's a remote DB through the public HTTP RESTful API)
 - Changes the replication checkpoint history of both DBs (source and remote) to mark that the source DB sequence number `SEQ_NUM` was processed
- After a crash (source or target):
 - The replicator compares the checkpoint history of both DBs and choses the highest common source DB sequence number `COMMON_SEQ_NUM`
 - Starts doing the above process but tells the `_changes` API to start from sequence number `COMMON_SEQ_NUM`

Replication Conflicts



Alice then decides to edit her version of the document:

In Alice's server	In Bob's server
<pre>{ "_id": "foo", "_rev": "2-72b1d3f01f345e7aa", "phone": "789123", "zip_code": "999" }</pre>	<pre>{ "_id": "foo", "_rev": "1-abcdef0123456789", "phone": "12345", "zip_code": "999" }</pre>

Not being aware of what Alice did to her copy, Bob decided to edit his copy:

In Alice's server	In Bob's server
<pre>{ "_id": "foo", "_rev": "2-72b1d3f01f345e7aa", "phone": "789123888", "zip_code": "999" }</pre>	<pre>{ "_id": "foo", "_rev": "2-111fff0893856def", "phone": "12345", "zip_code": "5984" }</pre>

Replication Conflicts



- Later on they decide to synchronize their DB's.
- What happens? Which version remains in DB?
- Both versions will remain. No data is lost on either side.
- But what will happen when :
 - GET <http://alice.com/contacts-db/foo>
 - GET <http://bob.com/contacts-db/foo>
- Simple: both request will return exactly the same document. When conflicting versions of a document are detected, CouchDB uses a deterministic algorithm to elect one version as the “winning” version. The elected “winning” version will be the same on all peers.

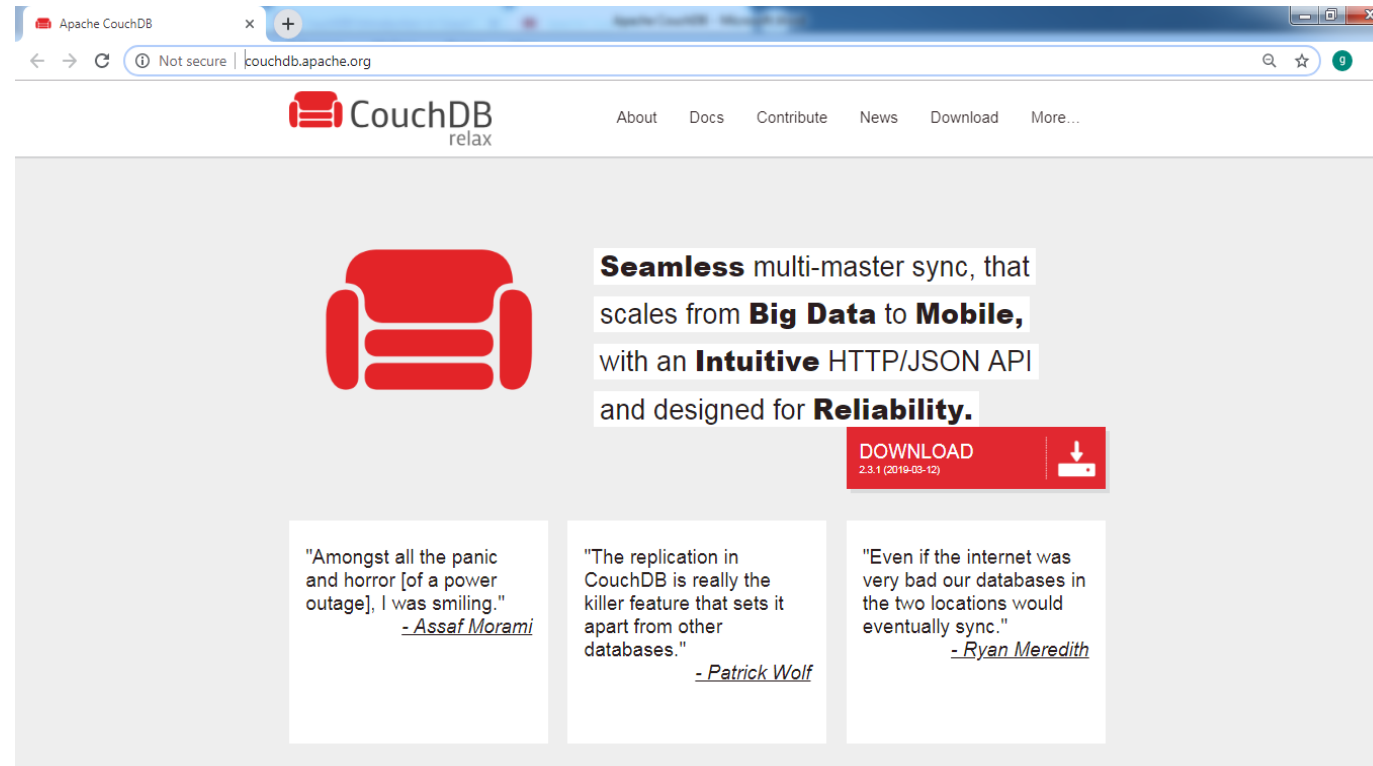


COUCHDB LAB

Installing Couch DB

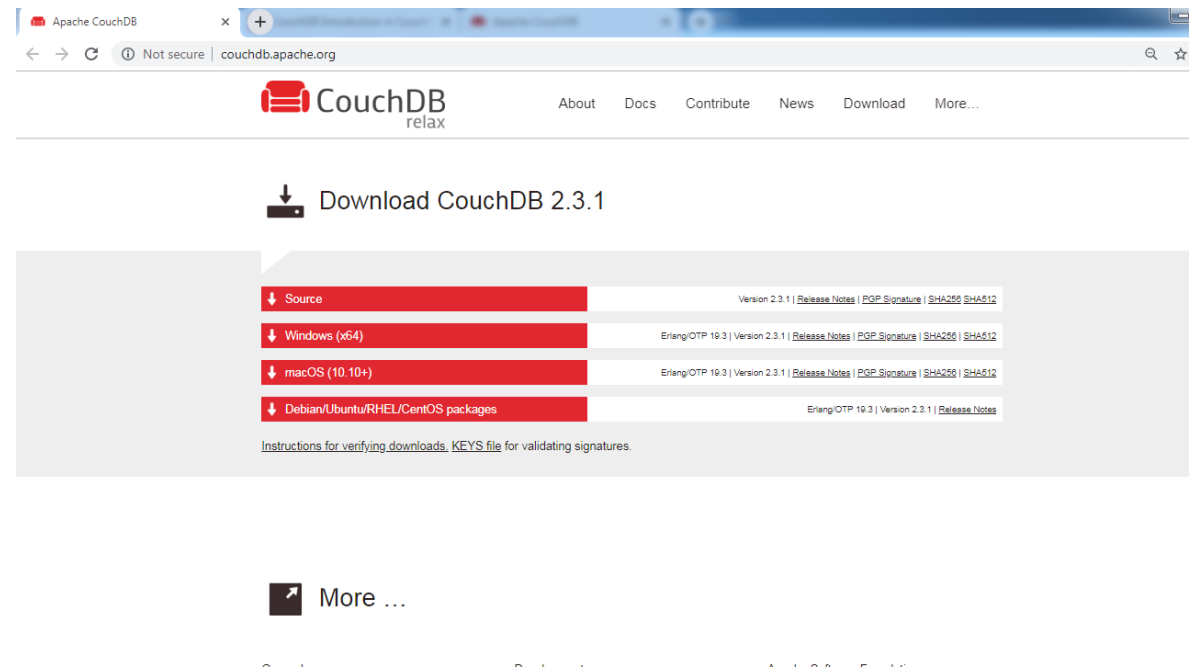


- For downloading the setup file of CouchDB, go to the official website: <http://couchdb.apache.org/>.



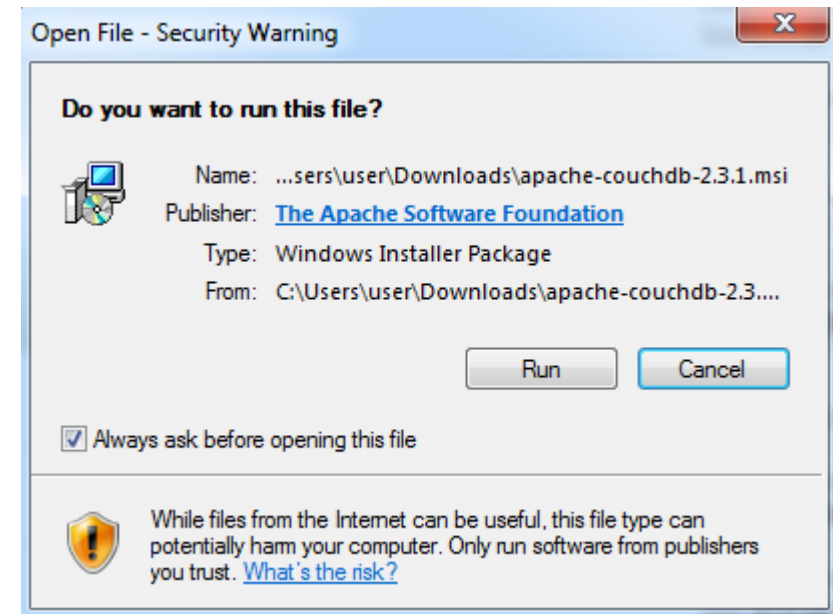


- Now, click on the Download option and then it leads to the page where various download links are provided.



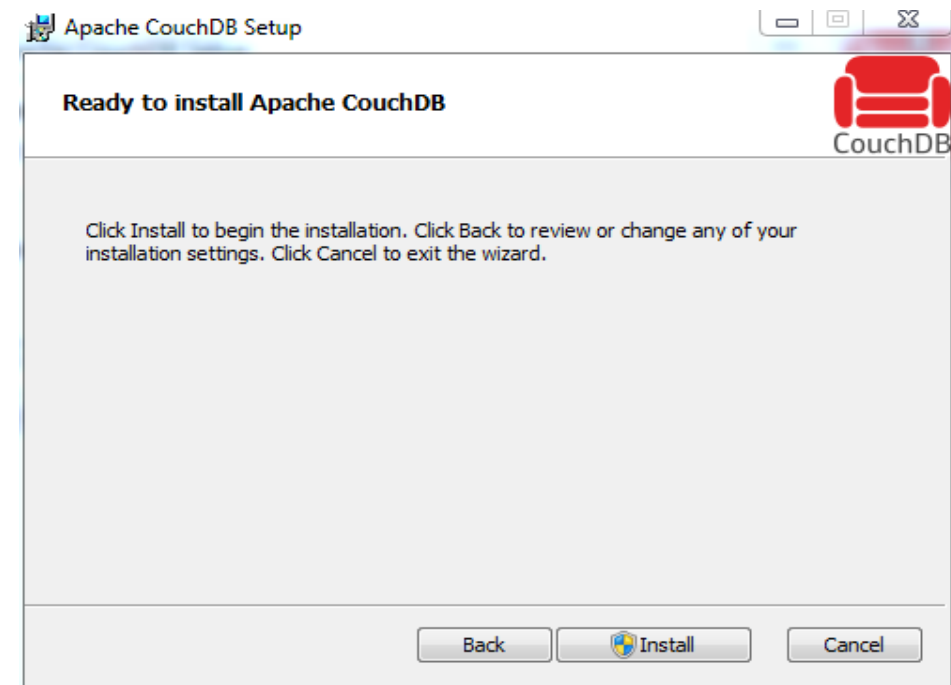


- Now, Click on the windows option. After 2-3 minutes, CouchDB will be download in your system in the form of Setup file.
- Now, run that setup file **apache-Couchdb-2.3.1.**






- After clicking on the run button, proceed with the installation. Default Settings are recommended, click on next button.
- Once you reach the Ready to install screen, click on Install. And after that, a confirmation dialog box will appear and in that click on 'yes' button. Your installation will complete in some seconds.






- After the complete installation opens the browser and type the following URL: <http://127.0.0.1:5984/> and open the link. If everything goes fine, the resulting output will appear:





- Now, by typing this URL: http://127.0.0.1:5984/_utils/, you can interact with the CouchDB web interface, which shows the index page of Futon as shown in the figure:





- Verify the CouchDB installation by going to the verify tab that shown in below figure. Now Click on the Verify installation.
- If your CouchDB is installed Successfully, a window will appear as shown below:

A screenshot of the CouchDB Verify Installation window. On the left is a dark sidebar with navigation links: Databases, Setup, Active Tasks, Configuration, Replication, Documentation, Verify (highlighted with a red background and a white checkmark icon), and Your Account. The main content area is titled "Verify CouchDB Installation" and contains a green "Verify Installation" button. Below the button is a table with two columns: "Test" and "Status". The table lists seven tests, all of which have a checkmark in the Status column, indicating successful installation.

Test	Status
Create Database	✓
Create Document	✓
Update Document	✓
Delete Document	✓
Create View	✓
Replication	✓

Futon: Create DB



Project Fauxton - 127.0.0.1:5984/_utils/#

Databases

Create Database {}JSON

Name	Size	# of Docs
big_data	8.9 KB	9
db_test	33.2 KB	0
nosqlclass	7.3 KB	10
test	1.0 KB	3
test_db	33.2 KB	0

Create Database

Name of database Create

Fauxton on Apache CouchDB v. 2.3.1

Log Out

Showing 1-5 of 5 databases. Databases per page 20 « 1 »

Futon: Delete DB



Project Fauxton - 127.0.0.1:5984/_utils/#

Databases

Name	Actions
big_data	[Icons]
db_test	[Icons]
nosqlclass	[Icons]
test	[Icons]
test_db	[Icons]

Confirm Deletion

Warning: This action will permanently delete `test_db`. To confirm the deletion of the database and all of the database's documents, you must enter the database's name.

Cancel Delete Database

Fauxton on CouchDB v 2.3.1 Log Out

Showing 1–5 of 5 databases. Databases per page 20

Futon: Creating a document

The screenshot shows the Apache CouchDB Futon web interface. The browser address bar displays the URL 127.0.0.1:5984/_utils/#database/nosqlclass/_all_docs. The left sidebar contains navigation links: All Documents, Run A Query with Mango, Permissions, Changes, and Design Documents. The main content area shows the 'nosqlclass' database with a 'New Document' button. A red circle highlights the 'Create Document' button in the top right corner. Below the button, a form is visible with fields for _id, name, address, age, and email. The bottom section shows a JSON document being created, with the following content:

```
1 {  
2   "_id": "857a553b90c7a1e3cc27ee32f000182c",  
3   "Name": "Sophia",  
4   "age": 30,  
5   "Designation": "Professor",  
6   "address": {  
7     "street": "5 main street",  
8     "city": "Toronto",  
9     "country": "Canada" }  
10 }  
11 }
```

CouchDB creates a new database document, assigning it a new id.
You can edit the value of the id and can assign your own value in the form of a string.

Futon: Updating a document



Project Fauxton - database/test X

127.0.0.1:5984/_utils/#database/test/e21489cca5cebdbfc705912953002769

test > e21489cca5cebdbfc705912953002769

Save Changes Cancel 2

Upload Attachment Clone Document Delete

```
1 {
2   "_id": "e21489cca5cebdbfc705912953002769",
3   "_rev": "2-aa2881afc5d6c253f57357deaaa1731d",
4   "name": "Alex Smith",
5   "age": 20,
6   "email": "Alex@gmail.com",
7   "address": {
8     "house_num": 23,
9     "street": "Bluegrass Ct",
10    "city": "Waterloo",
11    "province": "Ontario"
12  }
13 }
```

1

Fauxton on Apache CouchDB v. 2.3.1

Log Out

Creating View



Project Fauxton - database/test X

127.0.0.1:5984/_utils/#/database/test/new_view

test

- All Documents
- Run A Query with Mango
- Permissions
- Changes
- Design Documents

Add New

- New Doc
- New View
- Mango Indexes

New View

Design Document ?

New document _design/ newDesignDoc

Index name ?

new-view

Map function ?

```
1 function (doc) {  
2   emit(doc._id, 1);  
3 }
```

Reduce (optional) ?

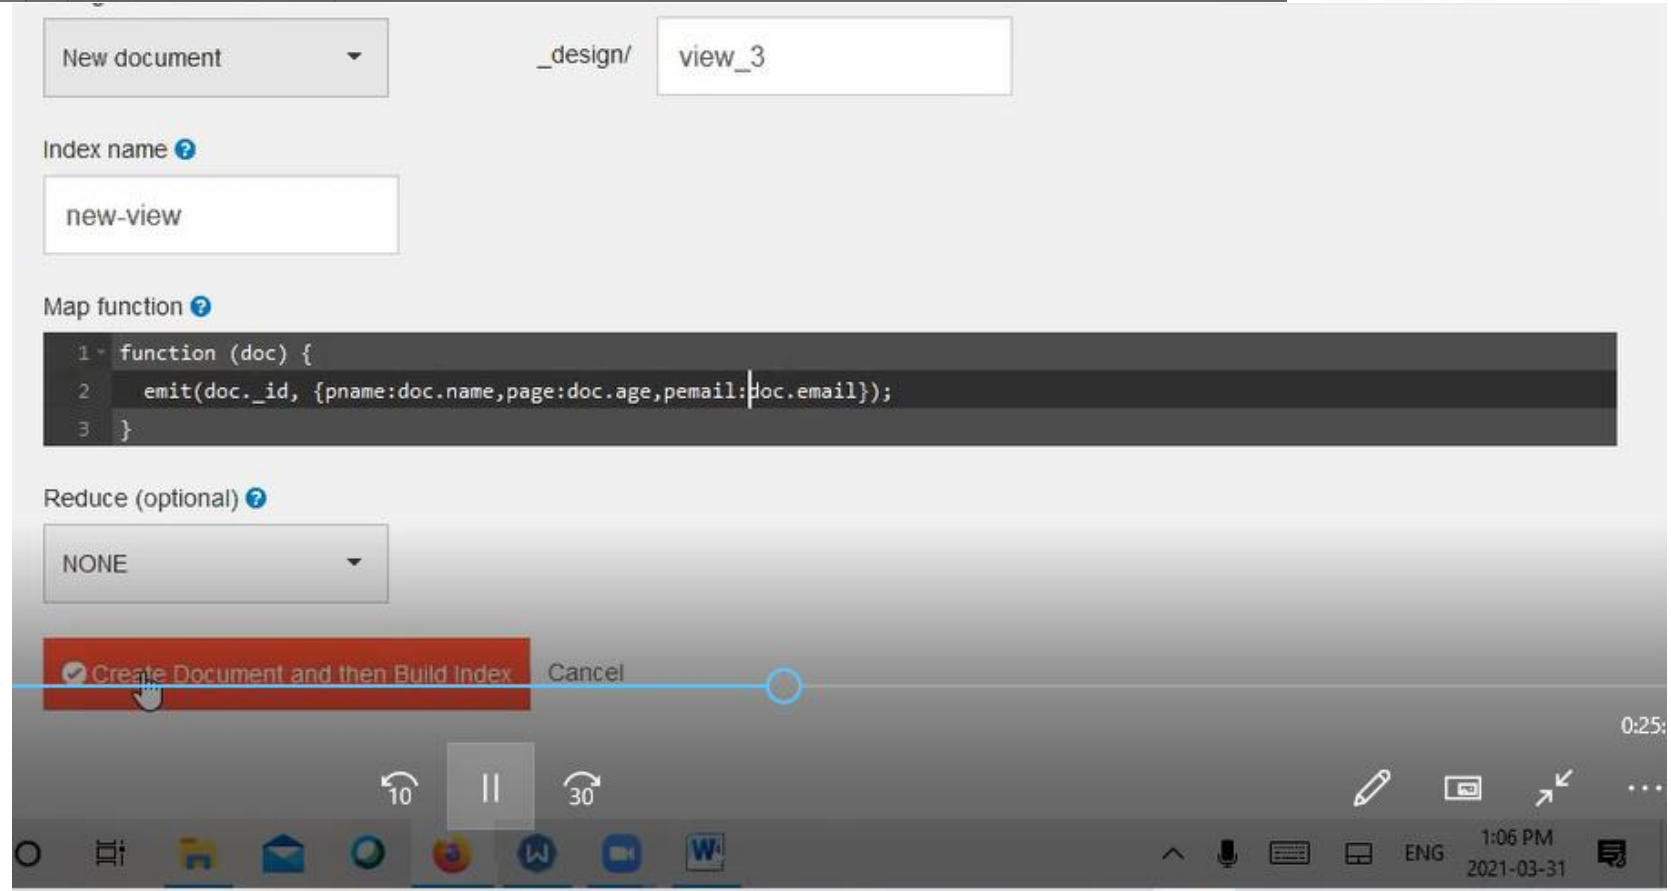
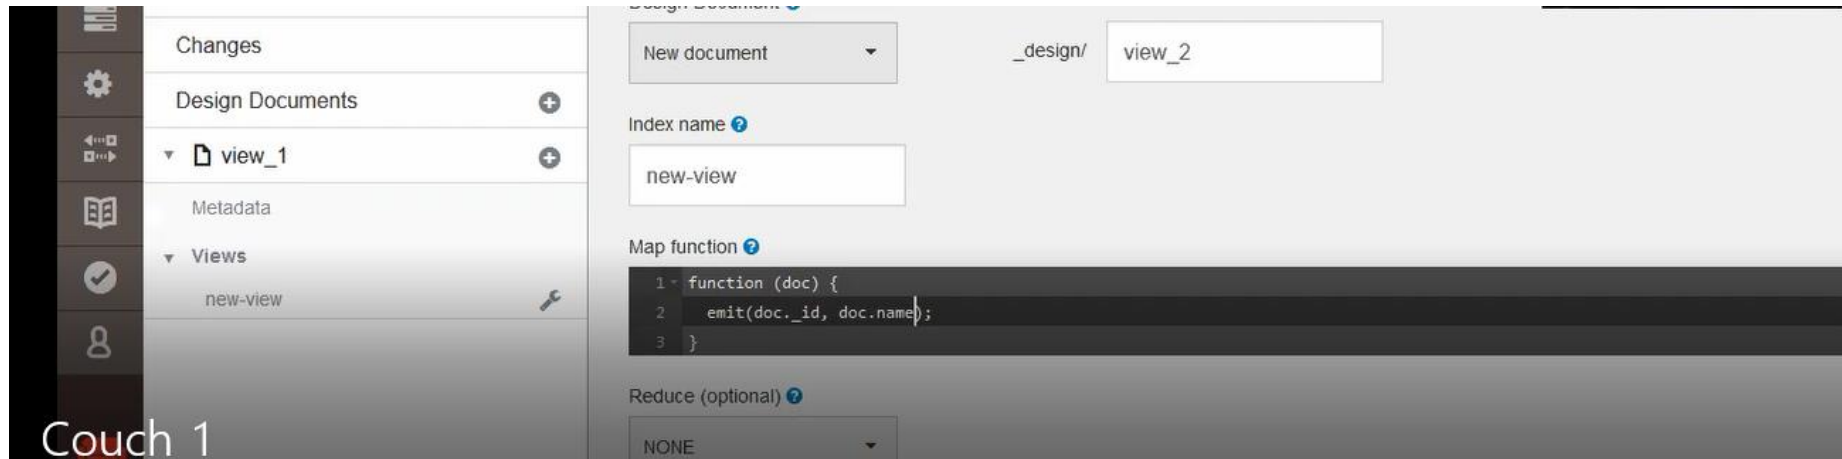
NONE

Create Document and then Build Index Cancel

Fauxton on Apache CouchDB v. 2.3.1

Log Out

127.0.0.1:5984/_utils/#/database/test/new_view



Command Line: Curl



- cURL utility is a way to communicate with CouchDB.
- The cURL utility is available in operating systems such as UNIX, Linux, Mac OS X and Windows. It is a command line utility using which user can access HTTP protocol straight away from the command line.
- To use Curl, goto command prompt , run it as administrator.
- Type `curl`
- You can access the homepage of the CouchDB by sending a GET request to the CouchDB instance installed.

```
curl http://127.0.0.1:5984/
```

Curl: Commands



- List databases: `curl -X GET http://127.0.0.1:5984/_all_dbs`
- Create database: `curl -X PUT http://127.0.0.1:5984/database_name`
- Verify database creation: `curl -X GET http://127.0.0.1:5984/all_dbs`
- Deleting database: `curl -X DELETE http://127.0.0.1:5984/database_name`
- Creating a document :

```
curl -X PUT http://127.0.0.1:5984/Nosql/"001" -d `
{ " Name " : " Sophia " ,
  " age " : " 30 " ,
  " Designation " : " Professor `,
  "address": {
    "street": "5 main street",
    "city": "Toronto",
    "country": "Canada"  }
} `
```

Curl: Commands



- Updating Documents:
 - First of all, get the revision id of the document that is to be updated. You can find the **_rev** of the document in the document itself.
 - Use revision id **_rev** from the document to update the document. Here we are updating the name from “Raju” to “Sophia Sandhu”.

```
curl -X GET http://127.0.0.1:5984/test/001 {"_id":"001",
  "_rev":"1-d75274b6a12d34cfc956b9a284f78f3b",
  "Name ":" Raju "}

curl -X PUT http://127.0.0.1:5984/test/001 -d"{
  \"_rev\": \"1-d75274b6a12d34cfc956b9a284f78f3b\",
  \"Name \": \"Sophia Sandhu\"}

{"ok":true,"id":"001","rev":"2-e28dfa41ac7493013449479db95444a5"}
```

Curl: Commands



- Deleting a Documents:
 - First of all, get the revision id of the document that is to be updated. You can find the **_rev** of the document in the document itself.
 - Use revision id **_rev** from the document to delete the document.

```
curl -X GET http://127.0.0.1:5984/test/001 {"_id": "001",  
  "_rev": "1-d75274b6a12d34cfc956b9a284f78f3b",  
  " Name ": " Raju "}
```

```
curl -X DELETE http://127.0.0.1:5984/test/001?_rev=1 d75274b6a12d34cfc956b9a284f78f3b
```


Future Reading



- Books:
 - CouchDB: The Definitive Guide, 2010, O'Reilly
Free online HTML version: <http://books.couchdb.org/relax>
 - CouchDB in Action, 2010, Manning
 - Beginning CouchDB, 2009, Apress
- Web:
 - Official site - <http://apache.couchdb.org>
 - Planet CouchDB - <http://planet.couchdb.org>
- Videos:
 - “CouchDB and me”, Damien Katz, <http://www.infoq.com/presentations/katz-couchdb-and-me>
 - Google Tech Talk, Chris Anderson, <http://www.youtube.com/watch?v=ESDBM9-U804>
 - SAPO Tech Talk, Jan Lehnardt, <http://developers.blogs.sapo.pt/14565.html>36

References

- Filipe David Manana 's Presentation.



Thanks