


Introduction to MongoDB

Sophia Sandhu



mongoDB

Agenda

- Mongo Db History
- MongoDB Characteristics
- MongoDB Objects
- Replication
- Sharding
- MongoDB Lab



MongoDB's History

- Developed by 10gen
 - Founded in 2007
 - A document-oriented, NoSQL database
 - Hash-based, schema-less database
 - No Data Definition Language - you can store hashes with any keys and values that you choose
 - Keys are stored as strings
 - Document Identifiers (`_id`) will be created for each document, field name reserved by system.
 - Uses BSON format
 - Based on JSON
 - Written in C++
 - Supports APIs (drivers) in many computer languages like : JavaScript, Python, Ruby, Perl, Java, C#, C++
-

What is MongoDB?

- Document-type NoSQL database
 - Open source (MongoDB Inc.)
 - Multi-platform
 - Documents in JSON-like format
 - JavaScript is integrated into the environment
 - Flexible schema
 - Editions:
 - MongoDB Community Server (Free)
 - MongoDB Enterprise Server (Commercial edition)
 - MongoDB Atlas (Cloud service)
 - Mongo shell used to manage databases
 - MongoDB Compass provides GUI-based database management
-

MongoDB Characteristics

- Dynamic schema
 - No DDL
 - Document-based database.
 - Query language via an API
 - Atomic writes and fully-consistent reads.
 - Support for embedded data models reduces I/O activity on database system.
 - Indexes support faster queries and can include keys from embedded documents and arrays
 - Master-slave replication with automated failover (replica sets)
 - Automated range-based partitioning of data (sharding).
 - Focuses on consistency and partition tolerance (CAP Theorem)
-

MongoDB Characteristics

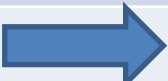

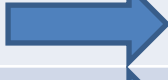
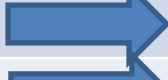

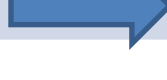
- MongoDB's replication facility, called [replica set](#), provides:
 - *automatic* failover
 - data redundancy.
- A [replica set](#) is a group of MongoDB servers that maintain the same data set, providing redundancy and increasing data availability.
- MongoDB provides horizontal scalability as part of its *core* functionality : [Sharding](#) distributes data across a cluster of machines.

MongoDB Objects

- Database – 0 or more databases
 - Collections - 1 or more Collections
 - Documents - 0 or more documents
 - Fields- 1 or more fields
-

RDBMS to NoSql



| RDBMS | | MongoDB |
|-----------|---|-------------------|
| Database | | Database |
| Table |  | Collection |
| Row |  | Document |
| Columns |  | Field |
| Index |  | Index |
| Join |  | Embedded document |
| Partition |  | Shard |

Schema Free Documents



- Unlike SQL databases, where you must determine and declare a table's schema before inserting data, MongoDB's collections, by default, do not require their documents to have the same schema. That is:
 - MongoDB does not need any pre-defined data schema.
 - Every document in a collection could have different data.
- A record in MongoDB is a document, which is a data structure composed of field and value pairs.
- MongoDB documents are similar to JSON objects. The values of fields may include other documents, arrays, and arrays of documents.

```
{
  "_id": {
    "$oid": "606ddf95a850f903fc25fce6"  },
  "FirstName": "Sophia",
  "LastName": "Sandhu",
  "major": "Computer Science",
  "phone":
    ["647-303-4567",
     "234-567-6789"],
  "gpa": 3.9
}
```

Schema Free Documents



Document 1

```
{
  "_id": { "$oid": "606ddf95a850f903fc25fce6" },
  "FirstName": "Sophia",
  "LastName": "Sandhu",
  "major": "Computer Science",
  "phone": [
    "647-303-4567",
    "234-567-6789"],
  "gpa": 3.5
}
```

Document 2

```
{
  "_id": { "$oid": "606de16349f3bd3718aa7361" },
  "StudentNo": "9900001",
  "FirstName": "John",
  "LastName": "Doe",
  "Grade": 85,
  "StudentType": "Domestic",
  "major": "Electrical"
}
```

- MongoDB does not need any pre-defined data schema.
- Every document in a collection could have different data.

Embedded Documents



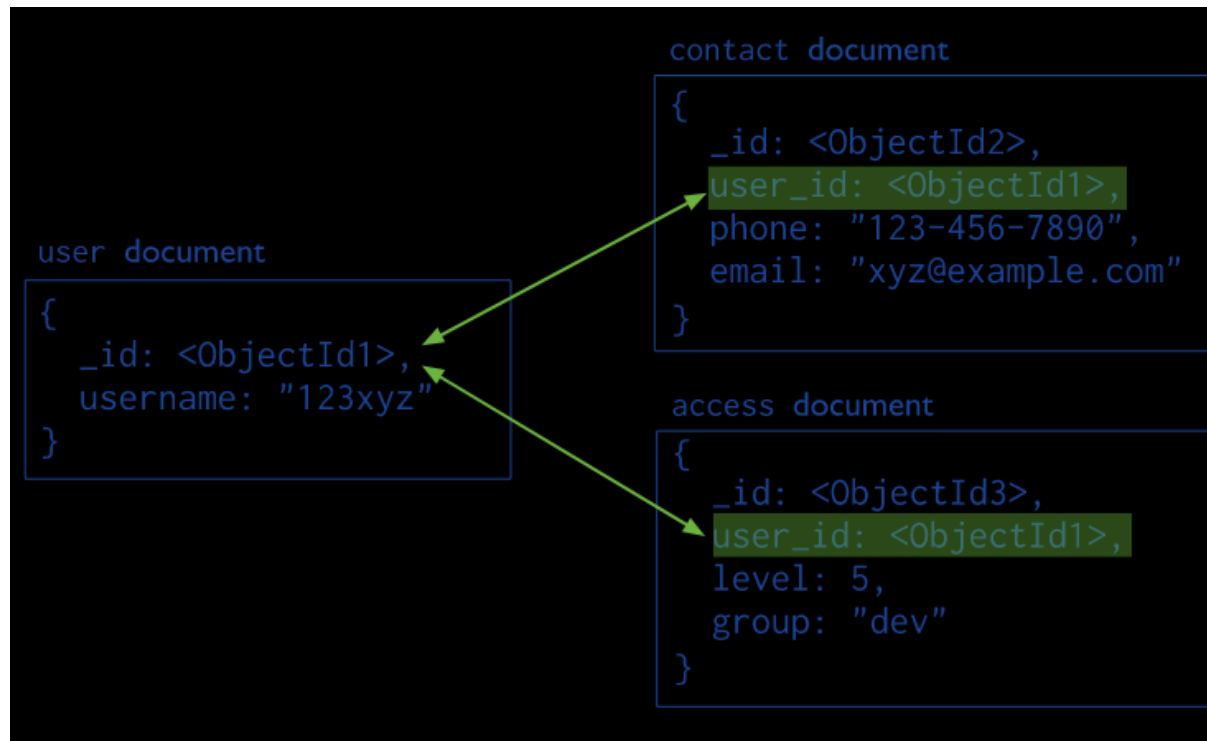
- Embedded documents capture relationships between data by storing related data in a single document structure.
- MongoDB documents make it possible to embed document structures in a field or array within a document.



References



- References store the relationships between data by including links or references from one document to another. Applications can resolve these references to access the related data.
- Data model using references to link documents.
- Both the ``contact`` document and the ``access`` document contain a reference to the ``user`` document.



Index



- B+ tree indexes
- An index is automatically created on the `_id` field (the primary key)
- Users can create other indexes to improve query performance or to enforce Unique values for a particular field.
 - Supports single field index as well as Compound index
 - Like SQL order of the fields in a compound index matters
 - If you index a field that holds an array value, MongoDB creates separate index entries for every element of the array
- Sparse property of an index ensures that the index only contain entries for documents that have the indexed field. (so ignore records that do not have the field defined)
- If an index is both unique and sparse – then the system will reject records that have a duplicate key value but allow records that do not have the indexed field defined

Replication




- A replica set in MongoDB is a group of mongod processes that maintain the same data set.
- Replica sets provide redundancy and high availability, and are the basis for all production deployments.
- With multiple copies of data on different database servers, replication provides a level of fault tolerance against the loss of a single database server.
- Replication can provide increased read capacity as clients can send read operations to different servers. Maintaining copies of data in different data centers can increase data locality and availability for distributed applications

MongoDB : Replication



- A replica set is a group of mongod instances that maintain the same data set.
- A replica set contains several data bearing nodes and optionally one arbiter node. Of the data bearing nodes, one and only one member is deemed the primary node, while the other nodes are deemed secondary nodes.
- The primary node receives all write operations. The primary records all changes to its data sets in its operation log, i.e. oplog.
- The secondaries replicate the primary's oplog and apply the operations to their data sets such that the secondaries' data sets reflect the primary's data set.
- If the primary is unavailable, an eligible secondary will hold an election to elect itself the new primary.



MongoDB : Replication



- When a primary does not communicate with the other members of the set for more than the configured `electionTimeoutMillis` period (10 seconds by default), an eligible secondary calls for an election to nominate itself as the new primary.
- The cluster attempts to complete the election of a new primary and resume normal operations.
- By default, clients read from the primary however, clients can specify a read preference to send read operations to secondary.

Sharding

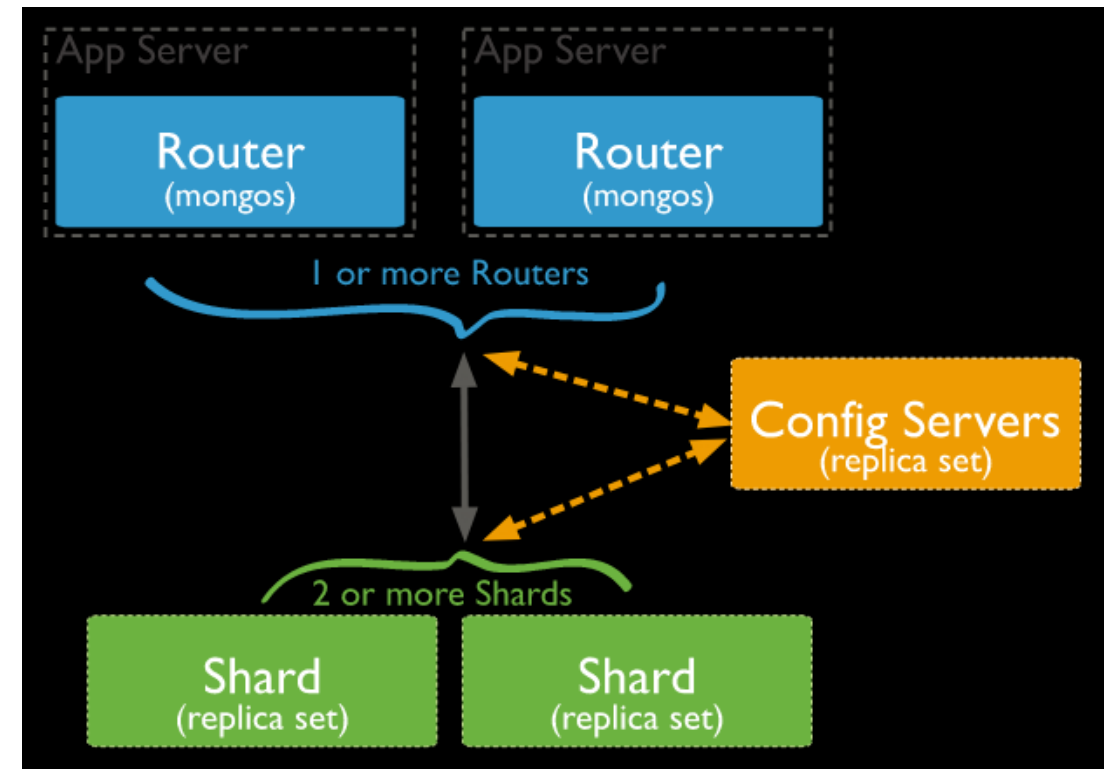


- Sharding is a method for distributing data across multiple machines. MongoDB uses sharding to support deployments with very large data sets and high throughput operations.
- Database systems with large data sets or high throughput applications can challenge the capacity of a single server. For example, high query rates can exhaust the CPU capacity of the server. Working set sizes larger than the system's RAM stress the I/O capacity of disk drives.
- There are two methods for addressing system growth: vertical and horizontal scaling.
 - **Vertical Scaling** involves increasing the capacity of a single server, such as using a more powerful CPU, adding more RAM, or increasing the amount of storage space.
 - Limitations in available technology may restrict a single machine from being sufficiently powerful for a given workload
 - **Horizontal Scaling** involves dividing the system dataset and load over multiple servers, adding additional servers to increase capacity as required. While the overall speed or capacity of a single machine may not be high, each machine handles a subset of the overall workload, potentially providing better efficiency than a single high-speed high-capacity server.

MongoDB: Sharding



- Mongo DB supports horizontal scaling through Sharding
- A MongoDB sharded cluster consists of the following components:
 - **Shard** : Each shard contains a subset of the sharded data. Each shard can be deployed as a replica set.
 - **mongos**: The mongos acts as a query router, providing an interface between client applications and the sharded cluster.
 - **config servers**: Config servers store metadata and configuration settings for the cluster.



MongoDB: Sharding



- MongoDB uses the shard keys to distribute the collection's documents across shards. The shard key consists of a field or multiple fields in the documents.
- A document's shard key value determines its distribution across the shards.
- To shard a populated collection, the collection must have an index that starts with the shard key.
- MongoDB partitions sharded data into chunks . Each chunk has an inclusive lower and exclusive upper range based on the shard key.
- MongoDB supports two sharding strategies for distributing data across sharded clusters:
 - Hashed Sharding
 - Ranged Sharding



MONGODB LAB

Installing MongoDB: Server



- Go to the MongoDB web site and download appropriate server (Community Server) in: <https://www.mongodb.com/download-center/community>
- Current version is 4.5.5
- Ensure the bin folder for Mongo is in the path to allow access in any directory
 - In bin folder (for Windows):
 - mongo.exe is the mongo shell
 - mongod.exe is the daemon
 - Similar executable file names are used for Linux and MacOS

Installing MongoDB

A screenshot of the MongoDB website's download center. The "Software" menu item in the top navigation bar is highlighted, and its dropdown menu is open, listing several options: "Community Server", "Enterprise Server", "Charts", "Compass", "Ops Manager", and "Connectors". Each option has a brief description. The background shows the "MongoDB Enterprise Server" section with a "download" button and a list of links like "Release notes", "Changelog", and "All version binaries".

mongoDB. Cloud **Software** Learn Solutions Docs

Search Contact Sign In [Try Free](#)

Select the server you want to download

- Community Server**
A free and open document database
- Enterprise Server**
Advanced features and security
- Charts**
Native visualization for MongoDB data
- Compass**
GUI for MongoDB
- Ops Manager**
On-prem management platform for MongoDB
- Connectors**
Easy integrations to your data estate

MongoDB Enterprise Server
ADVANCED FEATURES. PERFORMANCE GRADE.

[download](#)

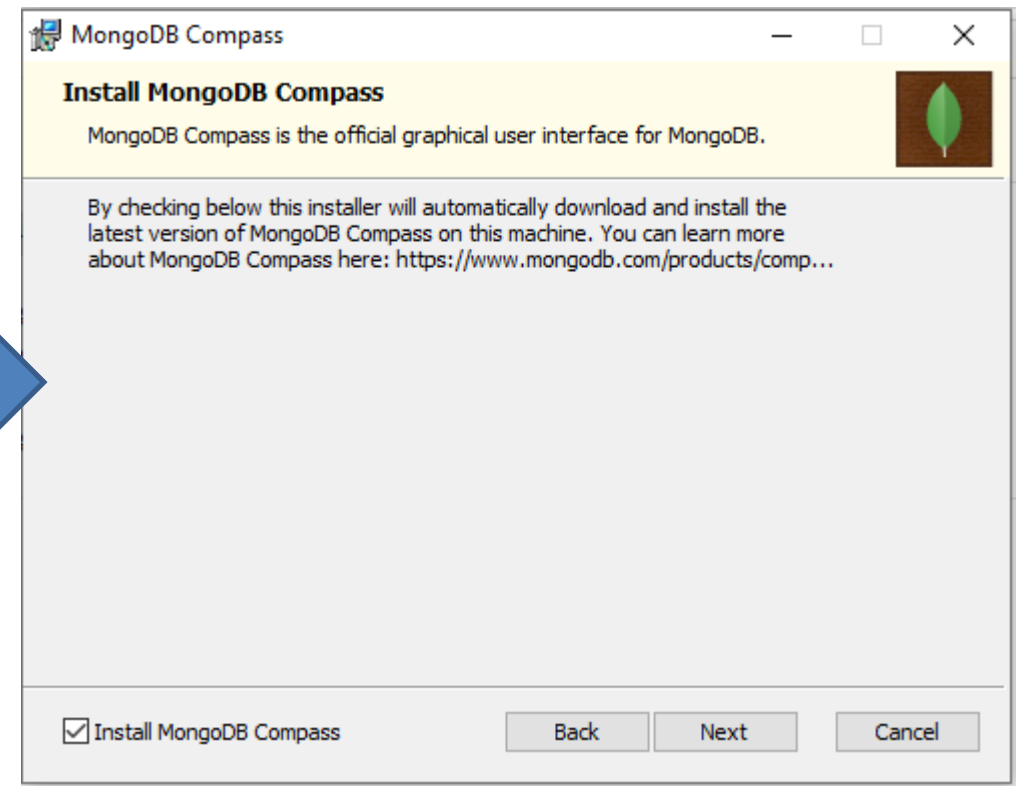
- Release notes
- Changelog
- All version binaries
- Installation instructions
- Download source (tgz)
- Download source (zip)

Version: 4.2.5 (current release)
Package: MSI
https://fastdl.mongodb.org/linux/mongodb-linux-x86_64-4.2.5.tgz

<https://www.mongodb.com/download-center/community>

Installing MongoDB

- Once download is complete open the msi file. Click Next in the start up screen.
- Accept the End-User License Agreement
- Click Next
- When you are on choose Setup type screen, make sure to select “Complete” option.
- Click next



Installing MongoDB



- Click on the Install button to start the installation.
- Follow the installation steps and click on “finish” button to complete the installation.

Using Mongo Shell



- Open the command prompt and goto the MongoDB Server->Server-> bin folder

- Enter mongo

```
cd C:\Program Files\MongoDB\Server\4.4\bin
```

```
C:\Program Files\MongoDB\
Server\4.4\bin>mongo
```

```
C:\Program Files\MongoDB\Server\4.4\bin>mongo
MongoDB shell version v4.4.2
connecting to: mongodb://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("4d89be2e-8d93-4582-9b05-97ecf6e28312") }
MongoDB server version: 4.4.2
---
The server generated these startup warnings when booting:
  2021-04-04T05:33:44.344-04:00: Access control is not enabled for the database. Read and write
configuration is unrestricted
---
---
  Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

  The monitoring data will be available on a MongoDB website with a unique URL accessible to y
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

  To enable free monitoring, run the following command: db.enableFreeMonitoring()
  To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---
>
```

MongoDB: Commands



- `show dbs` : display all databases
- `db`: display current database
- `use testDb` :
 - changes current db to testdb (testDb is the name of database)
 - You do not need to create the database before you switch.
 - MongoDB creates the database when you first store data in that database (such as create the first collection in the database).
- `db.myNewCollection2.insertOne({ x: 1 })`
- `db.myNewCollection3.createIndex({ y: 1 })`
- `db.createCollection('grade')`
 - Both the `insertOne()` and the `createIndex()` operations create their respective collection if they do not already exist
 - MongoDB provides the `db.createCollection()` method to explicitly create a collection with various options,

MongoDB: Commands



- `Db.getCollectionInfos()`
 - To retrieve the UUID for a collection
 - Collections are assigned an immutable UUID . The collection UUID remains the same across all members of a replica set and shards in a sharded cluster.

MongoDB: CRUD Operations



- CRUD operations *create, read, update, and delete documents*
- Create or insert operations add new documents to a collection.
- MongoDB provides the following methods to insert documents into a collection:
 - `db.collection.insertOne()`
 - `db.collection.insertMany()`

MongoDB: Insert Operation



```
db.inventory.insertOne(  
  { item: "canvas",  
    qty: 100,  
    tags: ["cotton"],  
    size: { h: 28, w: 35.5, uom: "cm" } }  
)
```

```
db.inventory.insertMany(  
  [  
    { item: "journal", qty: 25, tags: ["blank", "red"], size: { h: 14, w: 21, uom: "cm" } },  
    { item: "mat", qty: 85, tags: ["gray"], size: { h: 27.9, w: 35.5, uom: "cm" } },  
    { item: "mousepad", qty: 25, tags: ["gel", "blue"], size: { h: 19, w: 22.85 } }  
  ]  
)
```

MongoDB: Query Operation



- Assume the following 'Inventory' collection

```
[  
  { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },  
  { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },  
  { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },  
  { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },  
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" },  
  { item: "journal", qty: 25, tags: ["blank", "red"], dim_cm: [ 14, 21 ] },  
  { item: "notebook", qty: 50, tags: ["red", "blank"], dim_cm: [ 14, 21 ] },  
  { item: "paper", qty: 100, tags: ["red", "blank", "plain"], dim_cm: [ 14, 21 ] },  
  { item: "planner", qty: 75, tags: ["blank", "red"], dim_cm: [ 22.85, 30 ] },  
  { item: "postcard", qty: 45, tags: ["blue"], dim_cm: [ 10, 15.25 ] }  
]
```

MongoDB: Query Operation



- Use MongoDB Compass “filter” option or ‘mongo shell’ to query the collection
 - To display all documents use `{ }` in filter option or `db.Inventory.find()` in mongo shell
 - To display all documents where status is ‘A’ use : `{status : 'A'}` OR `db.Inventory.find({status : 'A'})`
 - To use ‘in’ operator use : `{ status: { $in: ["A", "D"] } }`
`db.Inventory.find({status: { $in: ["A", "D"] } })`
 - To use ‘and’ operator use : `{ status: "A", qty: { $lt: 30 } }`
OR `{ $and: [{ status: "A" }, { qty: { $lt: 30 } }] }`
`db.Inventory.find({$and: [{ status: "A" }, { qty: { $lt: 30 } }] })`
 - To use ‘or’ operator use : `{ $or: [{ status: "A" }, { qty: { $gt: 30 } }] }`
`db.Inventory.find({$or: [{ status: "A" }, { qty: { $lt: 30 } }] })`

MongoDB: Query Operation



- To use both 'or' and 'and' operator use :
`{item:'/^pa',{ $or: [{ status: "A" }, { qty: { $gt: 30 } }] }}`
- Querying embedded document : `{ size: { h: 14, w: 21, uom: "cm" } }`
- Querying nested field : `{ "size.uom":"cm" }`
- Querying an array : `{ tags: ["red", "blank"] }`
- Querying an array : `{ tags: { $all: ["red", "blank"] } }`
- To select few fields use : `db.inventory.find({ status: "A" }, { item: 1, status: 1 })`
`db.inventory.find({ status: "A" }, { status: 0, instock: 0 })`

MongoDB: Update Operation



```
db.students.insertMany([
  { _id: 1, test1: 95, test2: 92, test3: 90, modified: new Date("01/05/2020") },
  { _id: 2, test1: 98, test2: 100, test3: 102, modified: new Date("01/05/2020") },
  { _id: 3, test1: 95, test2: 110, modified: new Date("01/04/2020") },
```

- **`db.students.updateOne({ _id: 3 }, [{ $set: { "test3": 98, modified: "$$NOW"} }])`**

MongoDB: Update Operation



```
db.students3.insert([
  { "_id" : 1, "tests" : [ 95, 92, 90 ], "modified" : ISODate("2019-01-01T00:00:00Z") },
  { "_id" : 2, "tests" : [ 94, 88, 90 ], "modified" : ISODate("2019-01-01T00:00:00Z") },
  { "_id" : 3, "tests" : [ 70, 75, 82 ], "modified" : ISODate("2019-01-01T00:00:00Z") }
]);

db.students3.updateMany(
  { },
  [
    { $set: { average : { $trunc: [ { $avg: "$tests" }, 0 ] }, modified: "$$NOW" } },
    { $set: { grade: { $switch: {
      branches: [
        { case: { $gte: [ "$average", 90 ] }, then: "A" },
        { case: { $gte: [ "$average", 80 ] }, then: "B" },
        { case: { $gte: [ "$average", 70 ] }, then: "C" },
        { case: { $gte: [ "$average", 60 ] }, then: "D" }
      ],
      default: "F"
    } } } }
  ]
)
```

MongoDB: Delete Operation



- `db.Inventory.deleteOne({ "_id" : ObjectId("563237a41a4d68582c2509da") });`
- `db.Inventory.deleteOne({ "qty:" : { $gt:100} })`
- `db.Inventory.deleteMany({ "qty" : { $gt:100} })`
- `db.Inventory.remove({ qty: { $gt: 20 } })`

MongoDB: Aggregation Operation



- `db.orders.aggregate([
 { $match: { status: "A" } },
 { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }
])`
 - **First Stage:** The `$match` stage filters the documents by the `status` field and passes to the next stage those documents that have status equal to "A".
 - **Second Stage:** The `$group` stage groups the documents by the `cust_id` field to calculate the sum of the amount for each unique `cust_id`.
- `db.orders.distinct("cust_id")`

Views



- A MongoDB view is a queryable object whose contents are defined by an aggregation pipeline on other collections or views.
- MongoDB does not persist the view contents to disk.
- A view's content is computed on-demand when a client queries the view.
- MongoDB can require clients to have permission to query the view.
- MongoDB does not support write operations against views.

View



Using `db.CreateCollection()`

```
db.createCollection(  
    "<viewName>",  
    {  
        "viewOn" : "<source>",  
        "pipeline" : [<pipeline>],  
        "collation" : {  
<collation> }  
    }  
)
```

Using `db.CreateView()`

```
db.createView(  
    "<viewName>",  
    "<source>",  
    [<pipeline>],  
    {  
        "collation" : {  
<collation> }  
    }  
)
```

MONGO ATLAS

Mongo Atlas



The screenshot shows the MongoDB Atlas website. The navigation menu at the top includes links for Cloud, Software, Pricing, Learn, Solutions, and Docs. The 'Cloud' link is highlighted with a black circle. Below the navigation menu, the 'Atlas' section is visible, describing it as a 'Fully managed cloud database'. Other sections listed include 'Atlas Data Lake', 'Atlas Search', 'Realm', and 'Charts'. A green 'Start free' button is prominently displayed. On the right side of the page, there is a section for 'Cloud Provider & Region' which shows options for AWS, Google Cloud, and Azure. Below this, there is a table of regions categorized by North America, Europe, and Asia. The 'N. Virginia (us-east-1)' region is highlighted with a green box and labeled as 'FREE TIER AVAILABLE'.

Atlas
Fully managed cloud database

Atlas Data Lake
Query your AWS S3 and MongoDB data in-place

Atlas Search
Cloud-native full-text search engine

Realm
Application Development Services

Charts
Native visualization for MongoDB data

Start free

Already have an account? [Log in here](#) →

Cloud Provider & Region AWS, N. Virginia (us-east-1) ▾

aws Google Cloud Azure

Multi-Cloud, Multi-Region & Workload Isolation (M10+ clusters)
Distribute data across clouds or regions for improved availability and local read performance, or introduce replicas for workload isolation. [Learn more](#)

Create a **free tier cluster** by selecting a region with **FREE TIER AVAILABLE** and choosing the M0 cluster tier below.
★ Recommended region ⓘ

| NORTH AMERICA | EUROPE | ASIA |
|--|---|--|
| N. Virginia (us-east-1) ★ FREE TIER AVAILABLE | Stockholm (eu-north-1) ★ | Hong Kong (ap-east-1) ★ |
| Ohio (us-east-2) ★ | Ireland (eu-west-1) ★ FREE TIER AVAILABLE | Tokyo (ap-northeast-1) ★ |
| N. California (us-west-1) | London (eu-west-2) ★ | Seoul (ap-northeast-2) |
| Oregon (us-west-2) ★ FREE TIER AVAILABLE | Paris (eu-west-3) ★ | Singapore (ap-southeast-1) ★ FREE TIER AVAILABLE |

Mongo Atlas



← → ↻ 🏠 <https://www.mongodb.com/cloud/atlas> ... 🛡️ ☆

🔗 Nominate a project for the MongoDB Innovation Awards today! ➡

🌱 **mongoDB** Cloud Software Pricing Learn Solutions Docs 🔍 Contact Sign In [Try Free](#)

MongoDB Atlas

The most innovative cloud database service on the market, with unmatched data distribution and mobility across AWS, Azure, and Google Cloud, built in automation for resource and workload optimization, and so much more.

[Start free](#)

Already have an account? [Log in here](#) →

Next, Create an account

Cloud Provider & Region AWS, N. Virginia (us-east-1) ▾

Multi-Cloud, Multi-Region & Workload Isolation (M10+ clusters)
Distribute data across clouds or regions for improved availability and local read performance, or introduce replicas for workload isolation. [Learn more](#) ☐ OFF

Create a **free tier** cluster by selecting a region with **FREE TIER AVAILABLE** and choosing the **M0** cluster tier below.

★ Recommended region ⓘ

| NORTH AMERICA | EUROPE | ASIA |
|--|---|---|
| N. Virginia (us-east-1) ★ FREE TIER AVAILABLE | Stockholm (eu-north-1) ★ | Hong Kong (ap-east-1) ★ |
| Ohio (us-east-2) ★ | Ireland (eu-west-1) ★ FREE TIER AVAILABLE | Tokyo (ap-northeast-1) ★ |
| N. California (us-west-1) | London (eu-west-2) ★ | Seoul (ap-northeast-2) |
| Oregon (us-west-2) ★ FREE TIER AVAILABLE | Paris (eu-west-3) ★ | Singapore (ap-southeast-1) ★ FREE TIER AVAILABLE |
| Montreal (ca-central-1) | Frankfurt (eu-central-1) ★ FREE TIER AVAILABLE | Mumbai (ap-south-1) FREE TIER AVAILABLE |

Mongo Atlas: New User



- Click on Database Access to add new database user.

- 1: + ADD NEW DATABASE USER

2: Password Authentication fields

3: Select Atlas admin

4: Add User

Mongo Atlas: Network Access



- Click on Network Access to add IP Address.

■

A screenshot of the MongoDB Atlas web interface showing the 'Add IP Access List Entry' dialog box. The dialog box is white with a grey border and a close button (X) in the top right corner. It contains the following elements:

- Title:** 'Add IP Access List Entry'
- Text:** 'Atlas only allows client connections to a cluster from entries in the project's IP Access List. Each entry should either be a single IP address or a CIDR-notated range of addresses. [Learn more.](#)'
- Buttons:** Two buttons at the top: 'ADD CURRENT IP ADDRESS' (circled in blue) and 'ALLOW ACCESS FROM ANYWHERE'.
- Form Fields:** Two text input fields: 'Access List Entry:' with the value '208.78.41.139' and 'Comment:' with the placeholder text 'Optional comment describing this entry'.
- Footer:** A toggle switch labeled 'This entry is temporary and will be deleted in' followed by a dropdown menu set to '6 hours'. To the right are 'Cancel' and 'Confirm' buttons. The 'Confirm' button is circled in blue, and a large blue number '2' is written next to it.

The background of the screenshot shows the MongoDB Atlas interface with the 'Network Access' tab selected in the left sidebar. The URL in the browser's address bar is 'https://cloud.mongodb.com/v2/5fc70c456746f35a9291b9fc#security/network/accessList/addToAccessList'.

2

Mongo Atlas: Clusters


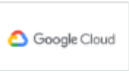






https://cloud.mongodb.com/v2/5fc70c456746f35a9291b9fc#clusters/edit?from=ctaClusterHeader 67%

Create New Cluster

Global Cluster Configuration >
















Cloud Provider & Region AWS, N. Virginia (us-east-1) ▾



 Multi-Cloud, Multi-Region & Workload Isolation (M10+ clusters)
Distribute data across clouds   or regions for improved availability and local read performance, or introduce replicas for workload isolation. [Learn more](#) ☐ OFF

Create a free tier cluster by selecting a region with **FREE TIER AVAILABLE** and choosing the M0 cluster tier below.

★ Recommended region ⓘ

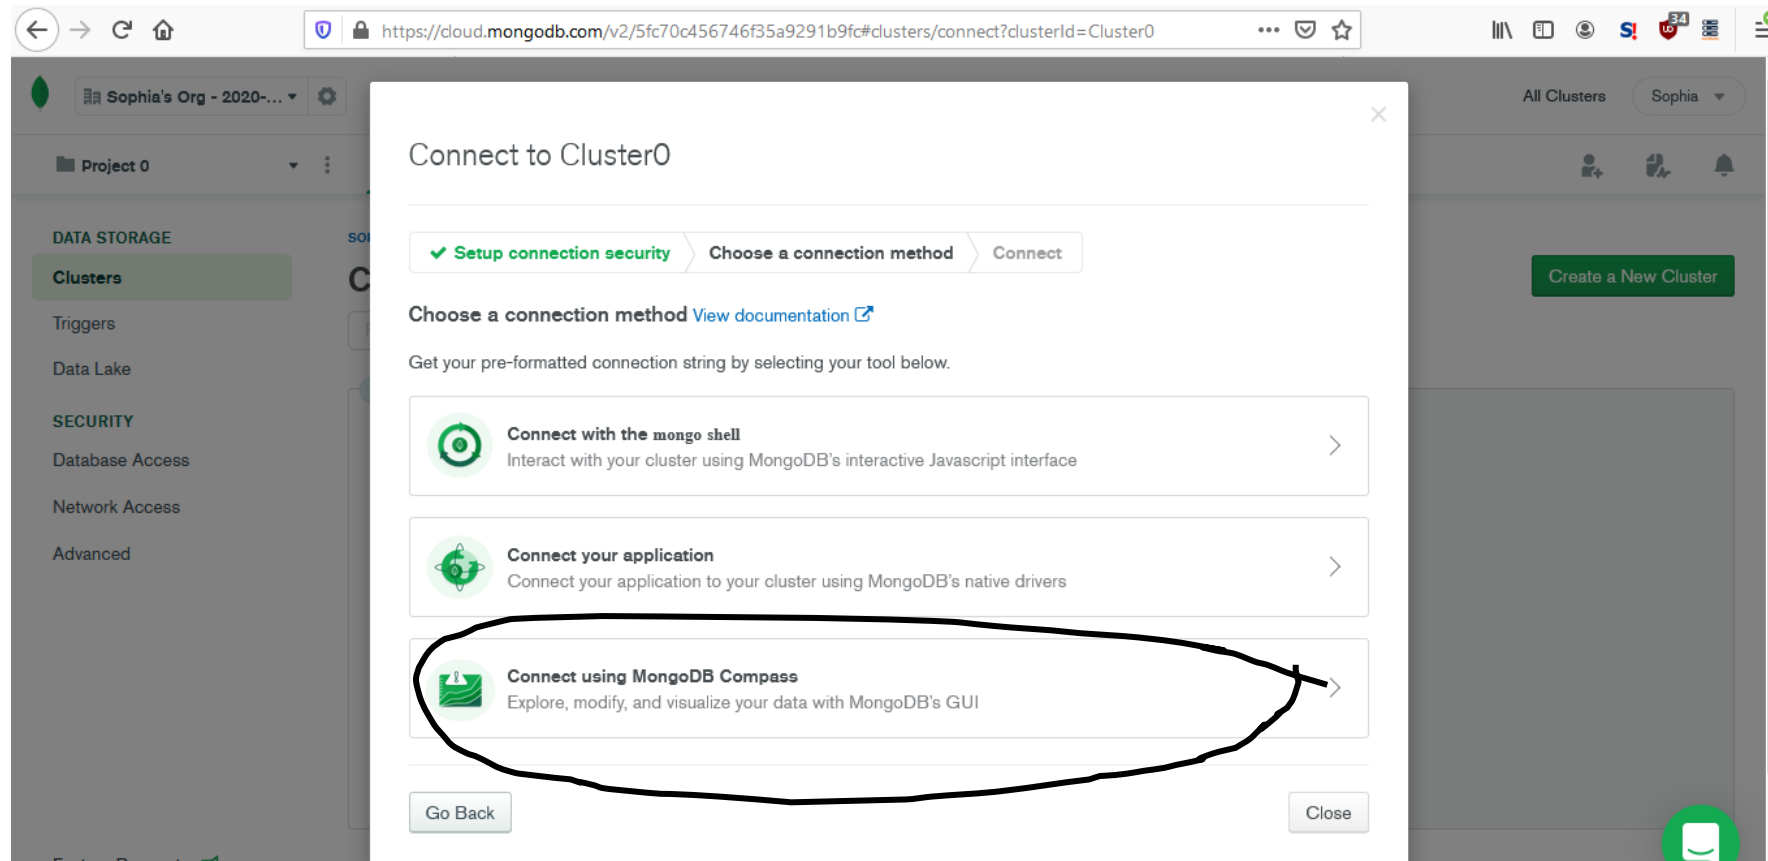
| NORTH AMERICA | EUROPE | ASIA |
|---|--|--|
|  N. Virginia (us-east-1) ★ FREE TIER AVAILABLE |  Stockholm (eu-north-1) ★ |  Hong Kong (ap-east-1) ★ |
|  Ohio (us-east-2) ★ |  Ireland (eu-west-1) ★ FREE TIER AVAILABLE |  Tokyo (ap-northeast-1) ★ |
|  N. California (us-west-1) |  London (eu-west-2) ★ |  Seoul (ap-northeast-2) |
|  Oregon (us-west-2) ★ FREE TIER AVAILABLE |  Paris (eu-west-3) ★ |  Singapore (ap-southeast-1) ★ FREE TIER AVAILABLE |
|  Montreal (ca-central-1) |  Frankfurt (eu-central-1) ★ FREE TIER AVAILABLE |  Mumbai (ap-south-1) FREE TIER AVAILABLE |

\$0.54/hour Pay-as-you-go! You will be billed hourly and can terminate your cluster anytime. Excludes variable data transfer, backup, and taxes.

Mongo Atlas



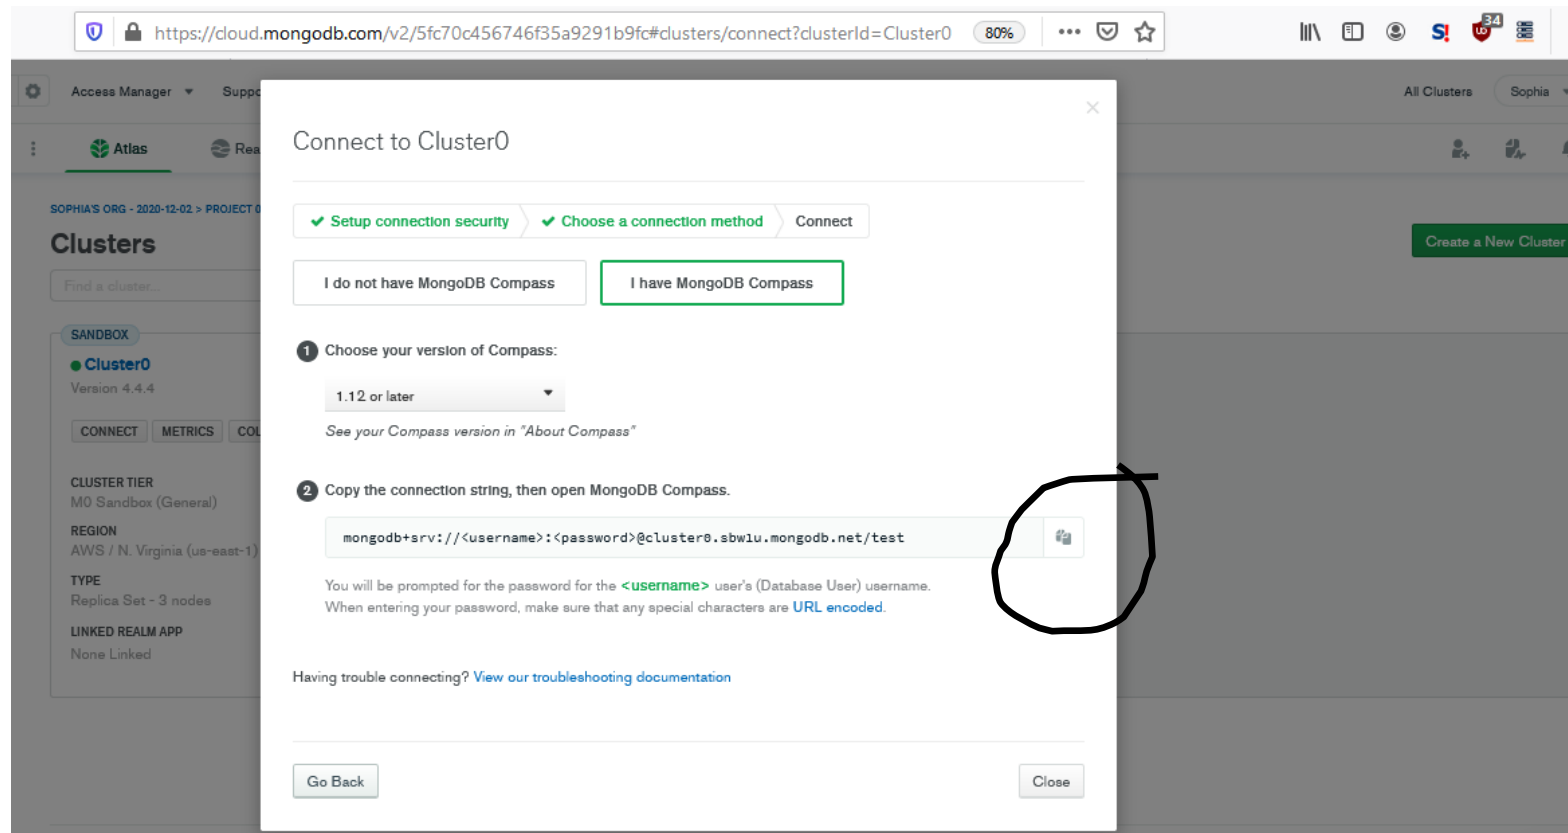
- Go to Clusters and click connect .
- Choose Connect using MongoDB Compass



Connection String



- Copy the connection string



MongoDB Compass

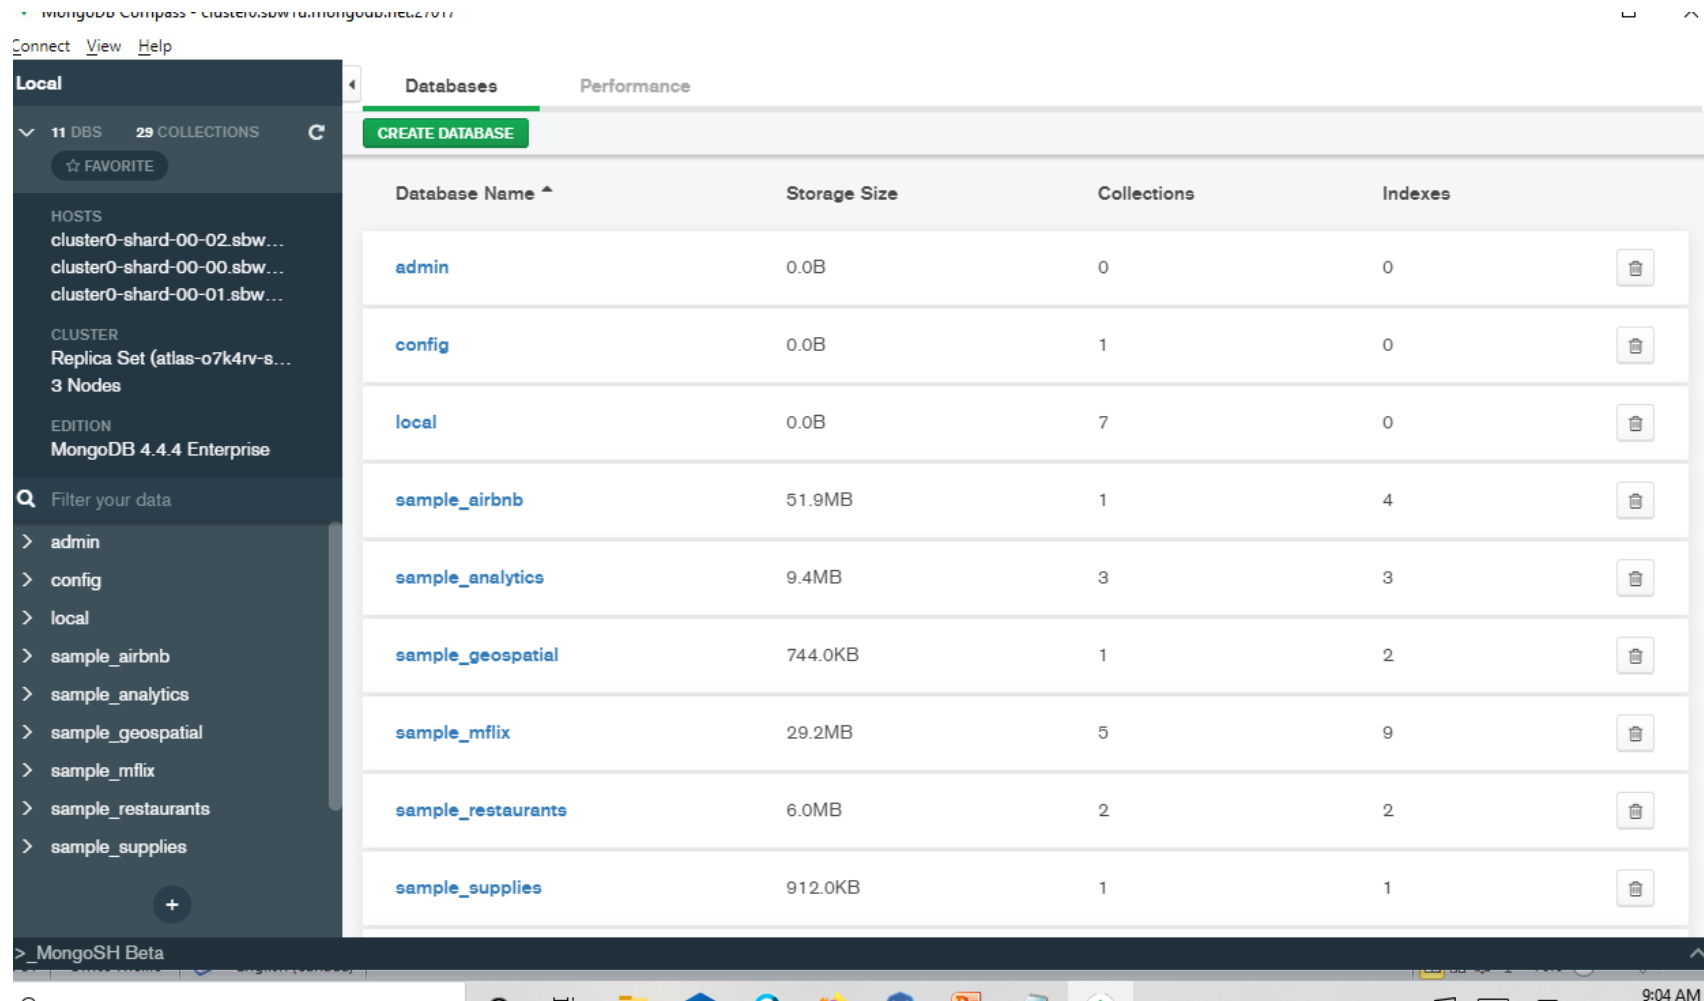


- Copy the connection string and paste it in MongoDB Compass under New Connection.

`mongodb+srv://<username>:<password>@cluster0.sbw1u.mongodb.net/test`

- Replace the username and password with actual values(username and password that we created under Database access)

MongoDB Compass



MongoDB Compass

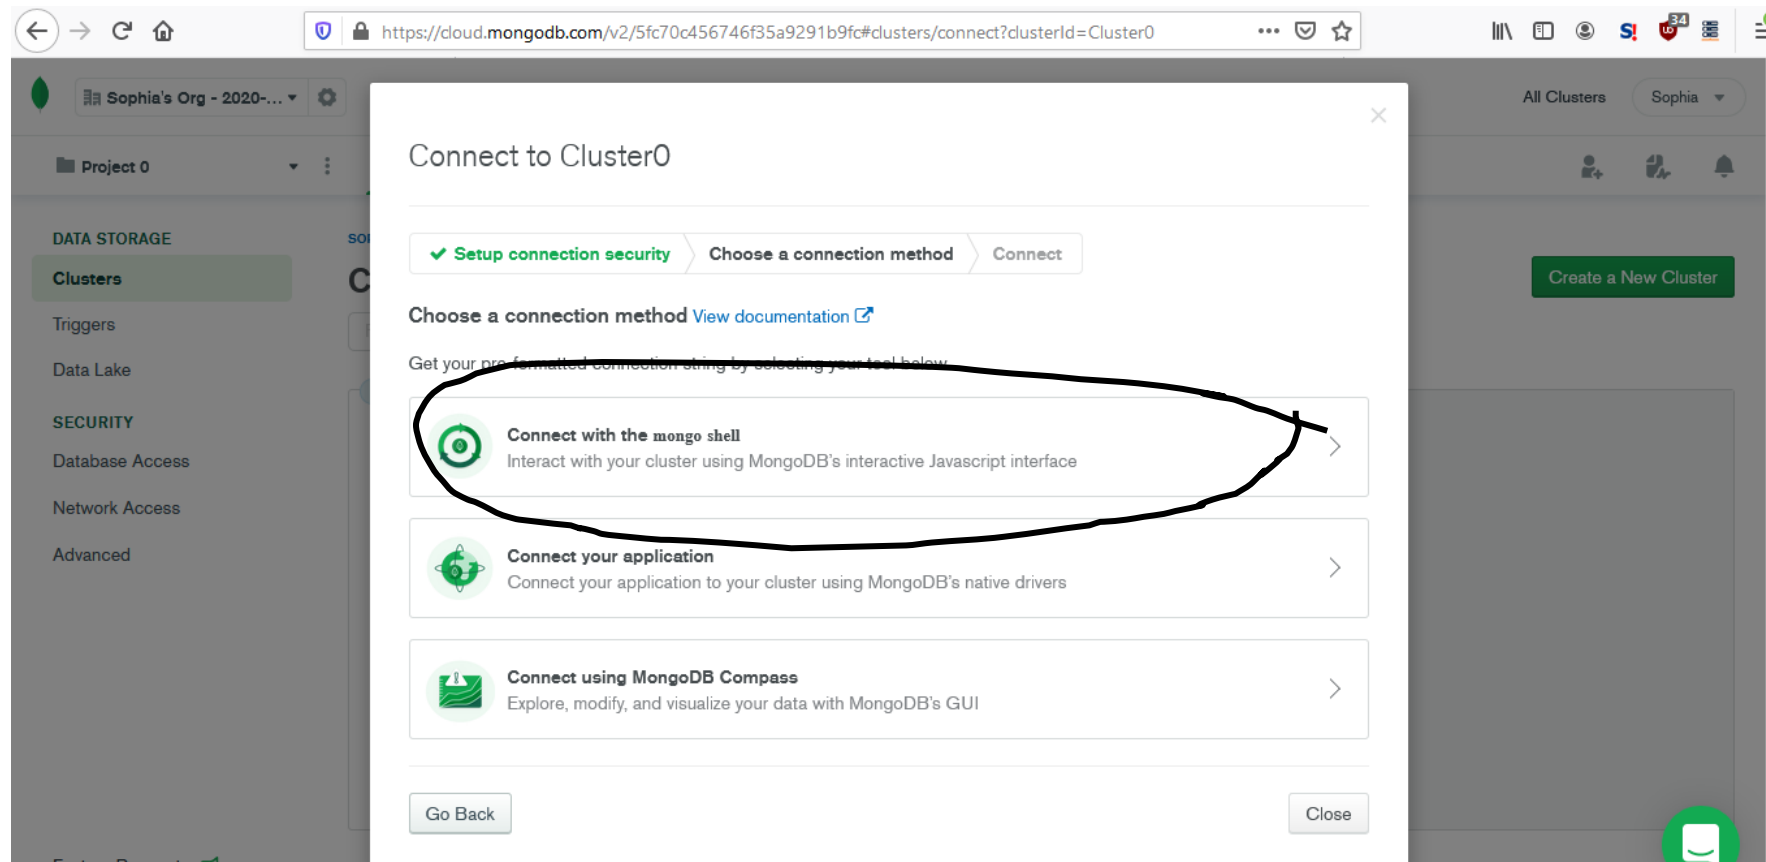


- Use any of the database to run Mongo queries on Mongo Collections.

Mongo Atlas



- Go to Clusters again and click connect .
- This time, choose Connect using Connect with the mongo shell



Mongo Atlas



- Copy the connection string

A screenshot of the MongoDB Atlas web interface. In the background, the 'Clusters' page for 'Project 0' is visible, showing details for 'Cluster0' (Sandbox, Version 4.4.4, Region: AWS / N. Virginia). Overlaid on this is a modal dialog titled 'Connect to Cluster0'. The dialog has a progress bar at the top with three steps: 'Setup connection security' (checked), 'Choose a connection method' (checked), and 'Connect'. Below this, there are two buttons: 'I do not have the mongo shell installed' and 'I have the mongo shell installed' (which is highlighted with a green border). The first step, 'Select your mongo shell version', shows a dropdown menu with '4.4' selected. The second step, 'Run your connection string in your command line', displays a text box with the connection string: `mongo "mongodb+srv://cluster0.sbw1u.mongodb.net/myFirstDatabase" --username <username>`. Below the text box, there is explanatory text about replacing 'myFirstDatabase' and encoding special characters. At the bottom of the dialog, there are 'Go Back' and 'Close' buttons, and a link to 'View our troubleshooting documentation'.

Mongo shell



- Open the command prompt and goto Mongo's bin folder.
- Paste the connection string and replace the username with actual <username>
- Run the mongo queries to see all databases that are available and work with any database/collection

```
C:\Program Files\MongoDB\Server\4.4\bin>mongo "mongodb+srv://cluster0.sbw1u.mongodb.net/myFirstDatabase" --username test
MongoDB shell version v4.4.2
Enter password:
connecting to: mongodb://cluster0-shard-00-00.sbw1u.mongodb.net:27017,cluster0-shard-00-02.sbw1u.mongodb.net:27017,cluster0-shard-00-01.sbw1u.mongodb.net:27017/myFirstD
atabase?authSource=admin&compressors=disabled&gssapiServiceName=mongodb&replicaSet=atlas-o7k4rv-shard-0&ssl=true
Implicit session: session { "id" : UUID("f0057d80-9339-4f4e-9ad2-248fbe22d6eb") }
MongoDB server version: 4.4.4
MongoDB Enterprise atlas-o7k4rv-shard-0:PRIMARY> _
```

Future Reading

- In Lynda.com, get the exercise files from
 - Learning MongoDB (Kirsten Hunter)
- MongoDB web site:
 - <https://www.mongodb.com>
 - MongoDB reference manual: •
<https://docs.mongodb.com/manual/reference/>
- Lynda.com course:
 - Learning MongoDB by Kirsten Hunter

References



- **MongoDB Documentation**
 - <https://docs.mongodb.com/manual/>

Thanks